

NAME : KRISH GUPTA

ROLL NO : 60 D15C

EXPERIMENT 1 : Phishing Website Detection using Logistic Regression

1. Dataset Source Dataset Name: **Phishing Website Detector**

Source: Kaggle

Link: <https://www.kaggle.com/datasets/eswarchandt/phishing-website-detector>

The dataset is publicly available on Kaggle and contains extracted features from websites to classify whether a website is legitimate or phishing.

2. Dataset Description Overview

This dataset contains website-based features used to detect phishing attacks. Each row represents one website instance.

Dataset Size

- Total Instances: 11,000+ (approx)
- Features: 30 input features
- Target Variable: Result

Target Variable

- Result = 1 → Legitimate Website
- Result = -1 → Phishing Website

(Binary classification problem)

Feature Description

The features are mostly extracted from:

- URL structure
- HTML & JavaScript content
- Domain-based features Examples of features:

- having_IP_Address • URL_Length
- Shortining_Service
- having_At_Symbol
- double_slash_redirecting
- Prefix_Suffix
- having_Sub_Domain
- SSLfinal_State
- Domain_registration_length
- HTTPS_token
- Request_URL
- URL_of_Anchor
- Links_in_tags
- SFH
- Submitting_to_email
- Abnormal_URL
- Redirect
- on_mouseover
- RightClick
- popUpWidnow
- Iframe
- age_of_domain
- DNSRecord
- web_traffic
- Page_Rank
- Google_Index
- Links_pointing_to_page
- Statistical_report

⚙ **Dataset Characteristics**

- All features are numeric (-1, 0, 1 format)
- No major missing values
- Balanced classification dataset

-
- Suitable for Logistic Regression

3. Mathematical Formulation of Logistic Regression

Logistic Regression is used for binary classification. **Step 1:**

Linear Combination

$$z = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$$

Where:

- w = weights
- x = feature values

Step 2: Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

This converts output into probability between 0 and 1.

Step 3: Decision Rule

If:

$$\sigma(z) \geq 0.5 \rightarrow \text{Class1}$$

Else:

$$\sigma(z) < 0.5 \rightarrow \text{Class0}$$

Step 4: Cost Function (Log Loss)

$$J(w) = -\frac{1}{m} \sum [y \log(h(x)) + (1 - y) \log(1 - h(x))]$$

Where:

- $h(x)$ = predicted probability
 - y = actual label
-

4. Algorithm Limitations

Logistic Regression has the following limitations:

1. Works only for linear decision boundaries.
2. Cannot capture complex non-linear relationships.
3. Sensitive to multicollinearity.
4. Performance decreases if dataset is highly imbalanced.
5. Outliers can affect decision boundary.

Not suitable when:

- Features have strong nonlinear dependency.
 - Very high-dimensional sparse datasets.
-

5. Methodology / Workflow Step 1: Data Collection

Dataset downloaded from Kaggle.

Step 2: Data Preprocessing

- Load dataset using Pandas
- Check null values
Separate features and target
- Train-test split (80%-20%)
- Feature scaling (if required)

Step 3: Model Training

- Apply Logistic Regression
- Fit model on training data

Step 4: Model Evaluation

- Predict on test data
 - Calculate Accuracy, Precision, Recall, F1-score
-

-

Workflow Diagram



6. Performance Analysis Evaluation

Metrics Used:

- Accuracy
- Precision
- Recall
- F1-Score
- Confusion Matrix

Example Results (Typical)

- Accuracy $\approx$ 94% – 97%
- Precision $\approx$ High
- Recall $\approx$ High
- F1-score $\approx$ Balanced Interpretation:
- High accuracy indicates strong classification ability.
- High recall means phishing websites are correctly detected.
- Low false negatives are important in cybersecurity applications.

7. Hyperparameter Tuning Hyperparameters

Tuned:

- C (Regularization strength)
- penalty (l1, l2)
- solver
- max_iter

Tuning Method Used:

GridSearchCV Example:

- Tested C values: [0.01, 0.1, 1, 10]
- Penalty: L1 and L2
- Solver: liblinear **Impact of Tuning:**

Before tuning:

- Accuracy $\approx$ 94%

After tuning:

- Accuracy improved to $\approx$ 96–97%
- Reduced overfitting

Better generalization

Code before tuning:

<pre>import pandas as pd import numpy as np import matplotlib.pyplot as plt import seaborn as sns import os from google.colab import files # Import files module from sklearn.model_selection import train_test_split from sklearn.linear_model import LogisticRegression, LinearRegression from sklearn.preprocessing import StandardScaler from sklearn.metrics import (accuracy_score, confusion_matrix,</pre>	<pre>classification_report, roc_curve, auc, mean_absolute_error, mean_squared_error, r2_score) print("Running both Logistic Regression and Linear Regression models...") # ===== === # 1. Load Dataset # ===== ===</pre>
---	---

```

file_path =
'/mnt/data/phishing.csv'
file_name_expected =
'phishing.csv'

df = None # Initialize df

if not os.path.exists(file_path): #
If the original path doesn't exist
    print(f"File '{file_path}' not
found. Attempting to load
'{file_name_expected}' from
current directory or prompting
for upload.")
    if
os.path.exists(file_name_expecte
d):
        print(f"Found
'{file_name_expected}' in current
directory.")
        df =
pd.read_csv(file_name_expected
)
    else:
        print(f"'{file_name_expected
}' not found in current directory.
Please upload
'{file_name_expected}'")
        uploaded = files.upload()
        if uploaded:
            uploaded_filename =
list(uploaded.keys())[0]
            if not
uploaded_filename.lower().ends
with('.csv'):
                raise
ValueError(f"Uploaded file
'{uploaded_filename}' is not a
CSV. Please upload a '.csv' file.")

        print(f"Reading uploaded
file: {uploaded_filename}")
        df =
pd.read_csv(uploaded_filename)

```

```

else:
    raise
FileNotFoundError(f"No file
uploaded. Please upload
'{file_name_expected}'")
else:
    print(f"Reading file from:
{file_path}")
    df = pd.read_csv(file_path)

if df is None:
    raise
FileNotFoundError("DataFrame
could not be loaded. Please
ensure a valid CSV file is available
or uploaded.")

print("Dataset Shape:", df.shape)
print(df.head())

#
=====
===
# 2. Data Preprocessing
#
=====
===

# Check for null values
print("\nMissing Values:\n",
df.isnull().sum())

# Assuming last column is target
(0 = Legitimate, 1 = Phishing)
# Note: Original dataset might
have -1 and 1 for target.
# For Linear Regression, we often
want 0 and 1, so let's normalize y
to 0/1 for both if it's -1/1.
# Let's check the unique values
of y to decide.
if -1 in df.iloc[:, -1].unique():

```

```

    print("\nConverting target
variable -1 to 0 for consistent
binary classification.")
    y_original = df.iloc[:, -
1].apply(lambda x: 1 if x == 1 else
0) # Convert -1 to 0
else:
    y_original = df.iloc[:, -1]

X = df.iloc[:, :-1]
y = y_original

#
=====
===
# 3. Train-Test Split
#
=====
===

X_train, X_test, y_train, y_test =
train_test_split(
    X, y, test_size=0.3,
    random_state=42)

#
=====
=====
#      Logistic Regression Model
#
=====
=====

print("\n=====
=====
==")
print("      LOGISTIC
REGRESSION MODEL      ")
print("=====
=====
==")

# 4. Feature Scaling (for Logistic
Regression)

```

```

scaler = StandardScaler()
X_train_scaled =
scaler.fit_transform(X_train)
X_test_scaled =
scaler.transform(X_test)

# 5. Logistic Regression Model
Training
logistic_model =
LogisticRegression(max_iter=100
00)
logistic_model.fit(X_train_scaled,
y_train)

# Predictions
y_pred_logistic =
logistic_model.predict(X_test_sc
aled)
y_prob_logistic =
logistic_model.predict_proba(X_t
est_scaled)[: ,1]

# 6. Model Evaluation
accuracy_logistic =
accuracy_score(y_test,
y_pred_logistic)
print("\nLogistic Regression
Model Accuracy:",
accuracy_logistic*100, "%")

print("\nLogistic Regression
Classification Report:\n")
print(classification_report(y_test,
y_pred_logistic))

# 7. Confusion Matrix
cm_logistic =
confusion_matrix(y_test,
y_pred_logistic)

plt.figure(figsize=(6,5))
sns.heatmap(cm_logistic,
annot=True, fmt='d',
cmap='Blues',

```



```

        xticklabels=['Legitimate',
'Phishing'],
        yticklabels=['Legitimate',
'Phishing'])
plt.title("Confusion Matrix –
Logistic Regression")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```

8. ROC Curve

```

fpr_logistic, tpr_logistic,
thresholds_logistic =
roc_curve(y_test,
y_prob_logistic)
roc_auc_logistic =
auc(fpr_logistic, tpr_logistic)

```

```

plt.figure(figsize=(7,6))
plt.plot(fpr_logistic, tpr_logistic,
label="ROC Curve (AUC = %0.4f)"
% roc_auc_logistic)
plt.plot([0,1], [0,1], linestyle='--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve – Logistic
Regression")
plt.legend()
plt.show()

```

9. Feature Importance

```

coefficients_logistic =
logistic_model.coef_[0]
feature_importance_logistic =
pd.Series(coefficients_logistic,
index=X.columns)
feature_importance_logistic =
feature_importance_logistic.sort
_values(ascending=False)

```

```

plt.figure(figsize=(10,6))
feature_importance_logistic.plot(
kind='bar')

```

```

plt.title("Feature Importance in
Logistic Regression")
plt.ylabel("Coefficient Value")
plt.xticks(rotation=90)
plt.show()

```

10. Mathematical Equation

```

print("\nLogistic Regression
Intercept (b0):",
logistic_model.intercept_[0])

```

```

print("Logistic Regression
Coefficients:")
for i, col in
enumerate(X.columns):
    print(f"{col}:
{logistic_model.coef_[0][i]}")

```

```

print("\nLogistic Regression
Equation:")
print("P(Phishing) = 1 / (1 + e^-
(b0 + b1x1 + b2x2 + ... + bnxn))")

```

```

#
=====
=====
#      Linear Regression Model
#
=====
=====

```

```

print("\n\n=====
=====
=====")
print("      LINEAR
REGRESSION MODEL      ")
print("=====
=====
=====")

```

4. Linear Regression Model

```

Training
linear_model =
LinearRegression()

```

```
linear_model.fit(X_train, y_train)
```

```
# Predictions (continuous values)
y_pred_cont_linear =
linear_model.predict(X_test)
```

```
# Convert to binary using
threshold 0.5 (for classification
evaluation)
y_pred_linear = [1 if val > 0.5 else
0 for val in y_pred_cont_linear]
```

```
# 5. Performance Metrics
mae_linear =
mean_absolute_error(y_test,
y_pred_cont_linear)
mse_linear =
mean_squared_error(y_test,
y_pred_cont_linear)
rmse_linear =
np.sqrt(mse_linear)
r2_linear = r2_score(y_test,
y_pred_cont_linear)
accuracy_linear =
accuracy_score(y_test,
y_pred_linear)
```

```
print("\n--- Linear Regression
Performance ---")
print("MAE:", mae_linear)
print("MSE:", mse_linear)
print("RMSE:", rmse_linear)
print("R2 Score:", r2_linear)
print("Classification Accuracy:",
accuracy_linear*100, "%")
```

```
# 6. Confusion Matrix
cm_linear =
confusion_matrix(y_test,
y_pred_linear)
```

```
plt.figure(figsize=(6,5))
```

```
sns.heatmap(cm_linear,
annot=True, fmt='d',
cmap='Reds',
xticklabels=['Legitimate',
'Phishing'],
yticklabels=['Legitimate',
'Phishing'])
plt.title("Confusion Matrix –
Linear Regression")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

```
# 7. Actual vs Predicted Graph
plt.figure(figsize=(7,6))
plt.scatter(y_test,
y_pred_cont_linear, alpha=0.4)
plt.axhline(0.5, color='red',
linestyle='--', label='Threshold
0.5')
plt.xlabel("Actual Label
(0=Legitimate, 1=Phishing)")
plt.ylabel("Predicted Continuous
Value")
plt.title("Actual vs Predicted –
Linear Regression")
plt.legend()
plt.show()
```

```
# 8. Feature Importance
coefficients_linear =
pd.Series(linear_model.coef_,
index=X.columns)
coefficients_linear =
coefficients_linear.sort_values(as
cending=False)
```

```
plt.figure(figsize=(10,6))
coefficients_linear.plot(kind='bar'
)
plt.title("Feature Importance –
Linear Regression")
plt.ylabel("Coefficient Value")
plt.xticks(rotation=90)
```

```
plt.show()
```

```
# 9. Regression Equation
print("\nLinear Regression
Intercept (b0):",
linear_model.intercept_)
```

```
print("Linear Regression
Coefficients:")
```

```
for i, col in
enumerate(X.columns):
    print(f"{col}:
{linear_model.coef_[i]}")

print("\nLinear Regression
Equation:")
print("y = b0 + b1x1 + b2x2 + ... +
bnxn")
```

Output

Running both Logistic Regression and Linear Regression models...
File '/mnt/data/phishing.csv' not found. Attempting to load 'phishing.csv' from current directory or prompting for upload.
'phishing.csv' not found in current directory. Please upload 'phishing.csv'.

[Choose files](#) No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving phishing.csv to phishing.csv

Reading uploaded file: phishing.csv

Dataset Shape: (11054, 32)

	Index	UsingIP	LongURL	ShortURL	Symbol@	Redirecting//	PrefixSuffix-	\
0	0	1	1	1	1	1	-1	
1	1	1	0	1	1	1	-1	
2	2	1	0	1	1	1	-1	
3	3	1	0	-1	1	1	-1	
4	4	-1	0	-1	1	-1	-1	

	SubDomains	HTTPS	DomainRegLen	...	UsingPopupWindow	IframeRedirection	\
0	0	1	-1	...	1	1	
1	-1	-1	-1	...	1	1	
2	-1	-1	1	...	1	1	
3	1	1	-1	...	-1	1	
4	1	1	-1	...	1	1	

	AgeofDomain	DNSRecording	WebsiteTraffic	PageRank	GoogleIndex	\
0	-1	-1	0	-1	1	
1	1	-1	1	-1	1	
2	-1	-1	1	-1	1	
3	-1	-1	0	-1	1	
4	1	1	1	-1	1	

	LinksPointingToPage	StatsReport	class
0	1	1	-1
1	0	-1	-1
2	-1	1	-1
3	1	1	1
4	-1	-1	1

[5 rows x 32 columns]

```

Missing Values:
  Index          0
UsingIP          0
LongURL          0
ShortURL         0
Symbol@         0
Redirecting//    0
PrefixSuffix-    0
SubDomains       0
HTTPS            0
DomainRegLen     0
Favicon          0
NonStdPort       0
HTTPSDomainURL   0
RequestURL       0
AnchorURL        0
LinksInScriptTags 0
ServerFormHandler 0
InfoEmail        0
AbnormalURL       0
WebsiteForwarding 0
StatusBarCust     0
DisableRightClick 0
UsingPopupWindow  0
IframeRedirection 0
AgeofDomain      0
DNSRecording      0
WebsiteTraffic    0
PageRank         0
GoogleIndex      0
LinksPointingToPage 0
StatsReport      0
class            0
dtype: int64

```

Converting target variable -1 to 0 for consistent binary classification.

```



=====
LOGISTIC REGRESSION MODEL
=====


```

Logistic Regression Model Accuracy: 92.67410310521555 %

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.93	0.90	0.91	1455
1	0.92	0.95	0.94	1862
accuracy			0.93	3317
macro avg	0.93	0.92	0.93	3317
weighted avg	0.93	0.93	0.93	3317





Logistic Regression Intercept (b0): 1.0330118259488312

Logistic Regression Coefficients:

Index: -0.13786576507727114

UsingIP: 0.612069322408425

LongURL: -0.1372353432133436

ShortURL: -0.5801489539777381

Symbol@: 0.18399100448047184

Redirecting//: 0.053924664953780475

PrefixSuffix-: 2.1074817443300997

SubDomains: 0.5304362803035259

HTTPS: 1.460013194497112

DomainRegLen: 0.036057361839474805

Favicon: -0.19501401005795177

NonStdPort: 0.3993004248504772

HTTPSDomainURL: -0.29890645848014313

RequestURL: 0.2709997750972327

AnchorURL: 2.289847464180018

LinksInScriptTags: 0.597681424631365

ServerFormHandler: 0.6983623700808214

InfoEmail: -0.30548316713082524

AbnormalURL: -0.06848615188404776

WebsiteForwarding: -0.41874823519745535

StatusBarCust: 0.16559360533909115

DisableRightClick: 0.02441602215632478

UsingPopupWindow: -0.12734143819098748

IframeRedirection: -0.12548919139443365

AgeofDomain: 0.07857792081575567

DNSRecording: 0.5401111735047753

WebsiteTraffic: 0.6347794707312321

PageRank: 0.14914382020914427

GoogleIndex: 0.46570623406860095

LinksPointingToPage: 0.44291386036724706

StatsReport: 0.22627308384756917

Logistic Regression Equation:

$P(\text{Phishing}) = 1 / (1 + e^{-(b_0 + b_1x_1 + b_2x_2 + \dots + b_nx_n)})$

LINEAR REGRESSION MODEL

--- Linear Regression Performance ---

MAE: 0.2090659724614079


MSE: 0.07454647434366536

RMSE: 0.27303200241668624


R2 Score: 0.6972561137914889


Classification Accuracy: 92.28218269520652 %

Confusion Matrix - Linear Regression



Actual vs Predicted - Linear Regression





Linear Regression Intercept (b0): 0.5525109708962723

Linear Regression Coefficients:

Index: -2.4136167860499328e-06

UsingIP: 0.04579682694709826

LongURL: -0.012666915428581792

ShortURL: -0.07387032010359144

Symbol@: 0.02774204479757832

Redirecting//: 0.026234483212919147

PrefixSuffix-: 0.05736175140251592

SubDomains: 0.0451847790628682

HTTPS: 0.20858704687554663

DomainRegLen: -0.005428443172240626

Favicon: -0.02115064245254335

NonStdPort: 0.03001967715616689

HTTPSDomainURL: -0.020122619011913073

RequestURL: 0.03012786632894973

AnchorURL: 0.24102135071331732

LinksInScriptTags: 0.0595473396675898

ServerFormHandler: 0.0648136003726666

InfoEmail: -0.019311631189308693

AbnormalURL: -0.02006073366643025

WebsiteForwarding: -0.0774611487626653

StatusBarCust: 0.016040351676942594

DisableRightClick: 0.000376277889236342

UsingPopupWindow: -0.001488103956969292

IframeRedirection: -0.017287116583704833

AgeofDomain: 0.001322202233390986

DNSRecording: 0.036382762946065406

WebsiteTraffic: 0.0564704443378479

PageRank: 0.023014553311820687

GoogleIndex: 0.04821837084017113

LinksPointingToPage: 0.05799601280732553

StatsReport: 0.03195593623427717

Linear Regression Equation:

$y = b_0 + b_1x_1 + b_2x_2 + \dots + b_nx_n$

Code after tuning:

```
# =====
# WITH HYPERPARAMETER TUNING
```



```
# =====
```

```
#  
=====
```

```
===  
# COMPLETE CODE (FINAL)
```

```
#  
=====
```

```
===  
  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import os
```

```
  
from google.colab import files  
from sklearn.model_selection  
import train_test_split,  
GridSearchCV  
from sklearn.linear_model  
import LogisticRegression, Ridge,  
Lasso  
from sklearn.preprocessing  
import StandardScaler  
from sklearn.metrics import (  
    accuracy_score,  
    confusion_matrix,  
    classification_report,  
    roc_curve, auc  
)
```

```
  
print("🌀 Running ML Models  
with Hyperparameter Tuning...")
```

```
#  
=====
```

```
===  
# 1. LOAD DATASET (AUTO FIX)
```

```
#  
=====
```

```
===
```

```
file_name = 'phishing.csv'
```

```
if not os.path.exists(file_name):
```

```
    print("📁 Upload your  
dataset:")  
    uploaded = files.upload()  
    file_name =  
list(uploaded.keys())[0]
```

```
df = pd.read_csv(file_name)
```

```
  
print("✅ Dataset Loaded  
Successfully")  
print("Shape:", df.shape)  
print(df.head())
```

```
#  
=====
```

```
===  
# 2. PREPROCESSING  
#
```

```
=====
```

```
===  
  
print("\nMissing Values:\n",  
df.isnull().sum())
```

```
  
# Convert target (-1 → 0)  
if -1 in df.iloc[:, -1].unique():  
    y = df.iloc[:, -1].apply(lambda  
x: 1 if x == -1 else 0)  
else:  
    y = df.iloc[:, -1]
```

```
X = df.iloc[:, :-1]
```

```
#  
=====
```

```
===  
# 3. TRAIN TEST SPLIT
```

```
#
=====

X_train, X_test, y_train, y_test =
train_test_split(
    X, y, test_size=0.3,
    random_state=42)

#
=====
=====
# LOGISTIC REGRESSION (TUNED)
#
=====
=====

print("\n===== LOGISTIC
REGRESSION (TUNED) =====")

scaler = StandardScaler()
X_train_scaled =
scaler.fit_transform(X_train)
X_test_scaled =
scaler.transform(X_test)

param_grid_log = {
    'C': [0.01, 0.1, 1, 10],
    'solver': ['liblinear', 'lbfgs']
}

grid_log =
GridSearchCV(LogisticRegression(
    max_iter=10000),
               param_grid_log,
    cv=5, scoring='accuracy')

grid_log.fit(X_train_scaled,
y_train)

best_log_model =
grid_log.best_estimator_
```

```
y_pred_log =
best_log_model.predict(X_test_s
caled)
y_prob_log =
best_log_model.predict_proba(X
_test_scaled)[:, 1]

print("Best Params:",
grid_log.best_params_)
print("Accuracy:",
accuracy_score(y_test,
y_pred_log) * 100)

print("\nClassification
Report:\n")
print(classification_report(y_test,
y_pred_log))

# Confusion Matrix
cm = confusion_matrix(y_test,
y_pred_log)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True,
fmt='d', cmap='Blues',
            xticklabels=['Legitimate',
'Phishing'],
            yticklabels=['Legitimate',
'Phishing'])
plt.title("Confusion Matrix –
Logistic Regression (Tuned)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# ROC Curve
fpr, tpr, _ = roc_curve(y_test,
y_prob_log)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(7,6))
plt.plot(fpr, tpr, label="AUC =
%0.4f" % roc_auc)
plt.plot([0,1], [0,1], linestyle='--')
plt.xlabel("False Positive Rate")
```

```

plt.ylabel("True Positive Rate")
plt.title("ROC Curve – Logistic
Regression (Tuned)")
plt.legend()
plt.show()

# Feature Importance
coef =
pd.Series(best_log_model.coef_[
0], index=X.columns)
coef.sort_values(ascending=False
).plot(kind='bar', figsize=(10,6))
plt.title("Feature Importance –
Logistic Regression")
plt.ylabel("Coefficient Value")
plt.xticks(rotation=90)
plt.show()

#
=====
=====
# RIDGE REGRESSION (TUNED)
#
=====
=====

print("\n===== RIDGE
REGRESSION (TUNED) =====")

param_grid_ridge = {'alpha':
[0.01, 0.1, 1, 10]}

grid_ridge =
GridSearchCV(Ridge(),
param_grid_ridge, cv=5)
grid_ridge.fit(X_train, y_train)

best_ridge =
grid_ridge.best_estimator_

y_pred_ridge =
best_ridge.predict(X_test)
y_pred_ridge_bin = [1 if i > 0.5
else 0 for i in y_pred_ridge]

```

```

print("Best Params:",
grid_ridge.best_params_)
print("Accuracy:",
accuracy_score(y_test,
y_pred_ridge_bin) * 100)

# Confusion Matrix
cm = confusion_matrix(y_test,
y_pred_ridge_bin)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True,
fmt='d', cmap='Reds',
xticklabels=['Legitimate',
'Phishing'],
yticklabels=['Legitimate',
'Phishing'])
plt.title("Confusion Matrix –
Ridge Regression")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Scatter Plot
plt.figure(figsize=(7,6))
plt.scatter(y_test, y_pred_ridge,
alpha=0.4)
plt.axhline(0.5, color='red',
linestyle='--')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted –
Ridge")
plt.show()

# Feature Importance
coef =
pd.Series(best_ridge.coef_,
index=X.columns)
coef.sort_values(ascending=False
).plot(kind='bar', figsize=(10,6))
plt.title("Feature Importance –
Ridge")
plt.xticks(rotation=90)

```

```

plt.show()

#
=====
=====
# LASSO REGRESSION (TUNED)
#
=====
=====

print("\n===== LASSO
REGRESSION (TUNED) =====")

param_grid_lasso = {'alpha':
[0.01, 0.1, 1, 10]}

grid_lasso =
GridSearchCV(Lasso(),
param_grid_lasso, cv=5)
grid_lasso.fit(X_train, y_train)

best_lasso =
grid_lasso.best_estimator_

y_pred_lasso =
best_lasso.predict(X_test)
y_pred_lasso_bin = [1 if i > 0.5
else 0 for i in y_pred_lasso]

print("Best Params:",
grid_lasso.best_params_)
print("Accuracy:",
accuracy_score(y_test,
y_pred_lasso_bin) * 100)

# Confusion Matrix

```

Output

```

cm = confusion_matrix(y_test,
y_pred_lasso_bin)
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True,
fmt='d', cmap='Greens',
xticklabels=['Legitimate',
'Phishing'],
yticklabels=['Legitimate',
'Phishing'])
plt.title("Confusion Matrix –
Lasso Regression")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Scatter Plot
plt.figure(figsize=(7,6))
plt.scatter(y_test, y_pred_lasso,
alpha=0.4)
plt.axhline(0.5, color='red',
linestyle='--')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("Actual vs Predicted –
Lasso")
plt.show()

# Feature Importance
coef =
pd.Series(best_lasso.coef_,
index=X.columns)
coef.sort_values(ascending=False
).plot(kind='bar', figsize=(10,6))
plt.title("Feature Importance –
Lasso")
plt.xticks(rotation=90)
plt.show()

```

Running ML Models with Hyperparameter Tuning...

Upload your dataset:

Choose files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving phishing.csv to phishing.csv

Dataset Loaded Successfully

Shape: (11054, 32)

	Index	UsingIP	LongURL	ShortURL	Symbol@	Redirecting//	PrefixSuffix-	\
0	0	1	1	1	1	1	1	-1
1	1	1	0	1	1	1	1	-1
2	2	1	0	1	1	1	1	-1
3	3	1	0	-1	1	1	1	-1
4	4	-1	0	-1	1	-1	-1	-1

	SubDomains	HTTPS	DomainRegLen	...	UsingPopupWindow	IframeRedirection	\
0	0	1	-1	...	1	1	1
1	-1	-1	-1	...	1	1	1
2	-1	-1	1	...	1	1	1
3	1	1	-1	...	-1	1	1
4	1	1	-1	...	1	1	1

	AgeofDomain	DNSRecording	WebsiteTraffic	PageRank	GoogleIndex	\
0	-1	-1	0	-1	1	1
1	1	-1	-1	1	-1	1
2	-1	-1	1	-1	1	1
3	-1	-1	0	-1	1	1
4	1	1	1	-1	1	1

	LinksPointingToPage	StatsReport	class
0	1	1	-1
1	0	-1	-1
2	-1	1	-1
3	1	1	1
4	-1	-1	1

[5 rows x 32 columns]

Missing Values:

Index	0
UsingIP	0
LongURL	0
ShortURL	0
Symbol@	0
Redirecting//	0
PrefixSuffix-	0
SubDomains	0
HTTPS	0

DomainRegLen	0
Favicon	0
NonStdPort	0
HTTPSDomainURL	0
RequestURL	0
AnchorURL	0
LinksInScriptTags	0
ServerFormHandler	0
InfoEmail	0
AbnormalURL	0
WebsiteForwarding	0
StatusBarCust	0
DisableRightClick	0
UsingPopupWindow	0
IframeRedirection	0
AgeofDomain	0
DNSRecording	0
WebsiteTraffic	0
PageRank	0
GoogleIndex	0
LinksPointingToPage	0
StatsReport	0
class	0

dtype: int64



===== LOGISTIC REGRESSION (TUNED) =====


Best Params: {'C': 0.1, 'solver': 'liblinear'}

Accuracy: 92.7645462767561


Classification Report:



	precision	recall	f1-score	support
0	0.93	0.90	0.92	1455
1	0.92	0.95	0.94	1862
accuracy			0.93	3317
macro avg	0.93	0.92	0.93	3317
weighted avg	0.93	0.93	0.93	3317





==== RIDGE REGRESSION (TUNED) =====
.. Best Params: {'alpha': 10}
Accuracy: 92.28218269520652





===== LASSO REGRESSION (TUNED) =====
Best Params: {'alpha': 0.01}
Accuracy: 91.86011456135061

